

Decentralized Group Messenger with Causally Ordered Multicast

Ahmet Uysal
auysal16@ku.edu.tr

Date: 15th of March

Contents

1	Introduction	2
1.1	Assumptions	2
2	Implementation	2
2.1	MessageDTO Struct	2
2.2	ConcurrentMessageSlice Struct	3
2.3	MessengerProcess Struct	3
2.3.1	PostMessage RPC Method	3
2.3.2	Helper Methods	5
2.4	Execution Flow of the Program	6
3	Deployment to AWS EC2 Instances	6
4	Possible Message Delivering Scenarios	7

1 Introduction

A decentralized unstructured peer-to-peer group messenger application is implemented based on Causally Ordered Multicasting algorithm using Go programming language and RPC protocol. AWS EC2 instances are used as the distributed platform.

1.1 Assumptions

- Multicast group size is static.
- Every peer has a file that contains addresses of all peers.
- Peers are assumed to function correctly once the initial connection is established. Therefore, the system does not implement any fault tolerance mechanisms.

2 Implementation

2.1 MessageDTO Struct

MessageDTO struct is how the messages are represented in my implementation. It consists of the identifier (IP address + port) of the sender, message timestamp, and the message string.

```

1 type MessageDTO struct {
2     Transcript string
3     OID        string
4     TimeStamp  []int
5 }

```

Code snippet 1: MessageDTO struct

MessageDTO struct also implements a method that checks whether a message should be delivered to application based on the Causally Ordered Multicasting algorithm. This method returns true if $ts(m)[sender] = VC_{receiver}[sender] + 1$ and $ts(m)[k] \leq VC_{receiver}[k]$ for all $k \neq sender$, and false otherwise.

```

1 func (message MessageDTO)
   shouldDeliverMessageToApplication(vectorClock []int,
   senderProcessIndex int) bool {
2     if message.TimeStamp[senderProcessIndex] !=
       vectorClock[senderProcessIndex]+1 {
3         return false
4     }
5     for k, time := range vectorClock {
6         if k == senderProcessIndex {
7             continue
8         }
9         if message.TimeStamp[k] > time {
10            return false
11        }
12    }
13    return true
14 }

```

Code snippet 2: shouldDeliverMessageToApplication Implementation

2.2 ConcurrentMessageSlice Struct

Decentralized group messenger needs to postpone the delivery of the messages that does not satisfy the requirements of `shouldDeliverMessageToApplication` method 2. In order to implement this, `MessengerProcess` structs need store queued messages in a data structure. However, Go Slices does not handle concurrency by themselves and mutexes are required to organize concurrent access and update scenarios. `ConcurrentMessageSlice` struct is created for this purpose. Implementation is based on a blog post by M. Nikolov. [1]

```

1 type ConcurrentMessageSlice struct {
2     sync.RWMutex
3     messages []MessageDTO
4 }
5 func (cms *ConcurrentMessageSlice) append(message MessageDTO) {
6     cms.Lock()
7     defer cms.Unlock()
8     cms.messages = append(cms.messages, message)
9 }
10 // Note: Removal of messages is implemented in MessengerProcess

```

Code snippet 3: ConcurrentMessageSlice Struct

2.3 MessengerProcess Struct

`MessengerProcess` struct is the main type that is responsible for delivery and interpretation of `MessageDTO` instances. It represents a peer process that is connected to Decentralized Group Messenger service. Each `MessengerProcess` stores the identifier (IP address + port) of the peer, current vector clock of the peer, queued messages that are waiting for to be delivered to the application, and a map that relates identifiers of the peers to indices of the vector clock.

```

1 type MessengerProcess struct {
2     OID string
3     QueuedMessages ConcurrentMessageSlice
4     VectorClock []int
5     VectorClockIndexMap map[string]int
6 }

```

Code snippet 4: MessengerProcess Struct

2.3.1 PostMessage RPC Method

`PostMessage` RPC method allows peers to transmit `MessageDTO` instances using RPC and handles received `MessageDTO` instances on peers. Go “net/rpc” package requires methods to follow a specific signature structure to be registered as an RPC call. In order to make a method publicly available to RPC calls, a method needs to:

- Have a type that is exported (`MessengerProcess` type is exported)
- Method itself is exported
- Method has two arguments, both exported (or builtin) types.
- The second argument is a pointer

- Method has the return type `error`

`PostMessage` is the only method of `MessengerProcess` type that satisfies all these requirements, and therefore it is can be registered. All other functions are helpers to implement message handling and does not need to be registered as RPC calls. RPC registration is implemented in the `main` function as shown below.

```

1 me := MessengerProcess{
2     OID: myAddress,
3     QueuedMessages: ConcurrentMessageSlice{
4         RWMutex: sync.RWMutex{},
5         messages: nil,
6     },
7     VectorClock:          vectorClock,
8     VectorClockIndexMap: peers,
9 }
10 _ = rpc.Register(&me)

```

Code snippet 5: RPC call registration

`PostMessage` method checks the received message for the requirements of being delivered to application. If the message satisfies the condition, it is delivered to application and the vector clock of the `MessengerProcess` is updated. Otherwise, the message is appended to the message queue. If a message is delivered to application and the vector clock is updated, `PostMessage` checks the message queue for messages that satisfy the delivery conditions based on the updated vector clocks. This procedure continues until there are no messages in the queue that satisfy delivery condition after a vector clock update. The `PostMessage` function implementation is given below along with `shouldDeliverMessageToApplication`, `deliverMessageToApplication`, and `popNextMessageToDeliver` helper methods.

```

1 func (messengerProcess *MessengerProcess) PostMessage(message
  MessageDTO, isSuccessful *bool) error {
2     if messengerProcess.shouldDeliverMessageToApplication(message){
3         messengerProcess.deliverMessageToApplication(message)
4         // if a message is delivered vector clock is updated
5         // we need to check queued messages
6         nextMessageToDeliver, shouldDeliver :=
            messengerProcess.popNextMessageToDeliver()
7         for shouldDeliver {
8             messengerProcess
9                 .deliverMessageToApplication(nextMessageToDeliver)
10            nextMessageToDeliver, shouldDeliver =
                messengerProcess.popNextMessageToDeliver()
11        }
12    } else {
13        // add message to queue
14        messengerProcess.QueuedMessages.append(message)
15    }
16    *isSuccessful = true
17    return nil
18 }

```

Code snippet 6: `PostMessage` RPC method

2.3.2 Helper Methods

`shouldDeliverMessageToApplication` is just a wrapper method around the `MessageDTO`'s `shouldDeliverMessageToApplication` method.

```

1 func (messengerProcess MessengerProcess)
2   shouldDeliverMessageToApplication(message MessageDTO) bool {
3     i := messengerProcess.VectorClockIndexMap[message.OID]
4     return message.shouldDeliverMessageToApplication(
5       messengerProcess.VectorClock, i)
6   }

```

Code snippet 7: `shouldDeliverMessageToApplication` helper method

`deliverMessageToApplication` updates the vector clock of the `MessengerProcess` and prints the received message to standard output.

```

1 func (messengerProcess *MessengerProcess)
2   deliverMessageToApplication(message MessageDTO) {
3   // update the vector clock
4   for index, time := range messengerProcess.VectorClock {
5     var maxTime int
6     if time > message.Timestamp[index] {
7       maxTime = time
8     } else {
9       maxTime = message.Timestamp[index]
10    }
11    messengerProcess.VectorClock[index] = maxTime
12  }
13  fmt.Printf("%s: %s\n", message.OID, message.Transcript)

```

Code snippet 8: `deliverMessageToApplication` helper method

`popNextMessageToDeliver` iterates over queued messages, pops and returns the first message that satisfies the delivery condition. It also returns a second flag value to indicate whether it found a message that satisfy the conditions.

```

1 func (messengerProcess *MessengerProcess) popNextMessageToDeliver() (MessageDTO, bool) {
2   var nextMessageToDeliver MessageDTO
3   var nextMessageToDeliverIndex int
4   messengerProcess.QueuedMessages.Lock()
5   defer messengerProcess.QueuedMessages.Unlock()
6   // iterate over queued messages until a match is found
7   for index, queuedMessage := range messengerProcess.QueuedMessages.messages {
8     if messengerProcess.shouldDeliverMessageToApplication(queuedMessage) {
9       nextMessageToDeliver = queuedMessage
10      nextMessageToDeliverIndex = index
11      break
12    }
13  }
14  // no message satisfies the condition
15  if nextMessageToDeliver.OID == "" {
16    return MessageDTO{}, false
17  }
18  messengerProcess.QueuedMessages.messages =
19    append(messengerProcess.QueuedMessages.messages[:nextMessageToDeliverIndex],
20      messengerProcess.QueuedMessages.messages[nextMessageToDeliverIndex+1:]...)
21  return nextMessageToDeliver, true

```

Code snippet 9: `popNextMessageToDeliver` helper method

2.4 Execution Flow of the Program

The program starts by getting its own IP address and reads the peers.txt file. It gives an error and exits if it couldn't find its own identifier inside this file. Every **MessengerProcess** starts with a vector clocks consisting of all zeros. Then, it registers the **PostMessage** method to RPC as explained in the previous sections. A goroutine is created to continuously listen for incoming RPC connections. Each **MessengerProcess** then waits until all peers listed inside peers.txt file is available by continuously trying to connect to each peer until it responds. After the connection is established, it informs the user and starts to accept message inputs. Vector clock of the **MessengerProcess** is increased by one only on message send events. Messages are multicasted by RPC calls to each peer.

3 Deployment to AWS EC2 Instances

Decentralized Group Messenger implementation is tested with five peers using AWS EC2 instances. I used PuTTY to transfer source files and connect to each client over SSH connection.

Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS (IPv4)	IPv4 Public IP
i-0097952516c871711	t2.micro	us-east-1b	running	2/2 checks ...	None	ec2-52-206-76-106.co...	52.206.76.106
i-04b332a0d2eed04...	t2.micro	us-east-1b	running	2/2 checks ...	None	ec2-54-172-5-91.comp...	54.172.5.91
i-05132f1a14882b8f	t2.micro	us-east-1b	running	2/2 checks ...	None	ec2-18-208-107-40.co...	18.208.107.40
i-06245518336700e...	t2.micro	us-east-1b	running	2/2 checks ...	None	ec2-54-234-28-3.comp...	54.234.28.3
i-06fb2213cef410e7e	t2.micro	us-east-1b	running	2/2 checks ...	None	ec2-3-87-4-159.comput...	3.87.4.159

Figure 1: Five AWS EC2 instances that are used in the deployment phase.

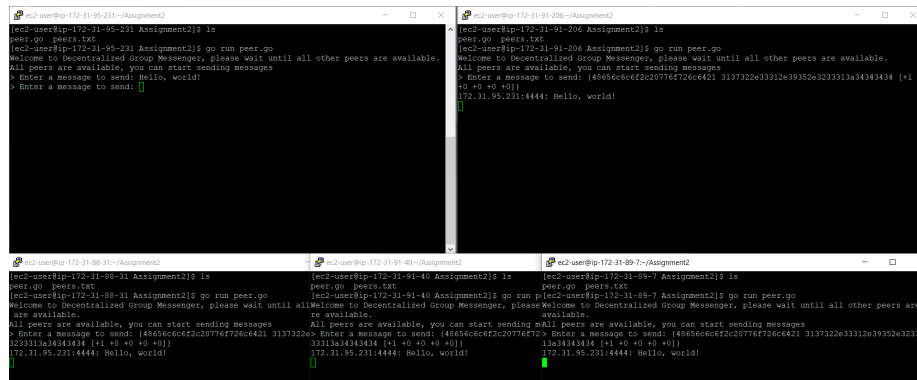


Figure 2: Initial connection and messaging between group of five peers running on AWS EC2 instances

In order to illustrate different scenarios for message delivery, I added a random delay between 0-5 seconds to each message transfer in `postMessageToPeer` function. Since this function is called in separate goroutines for each peer, sender will send the messages to each peer in a random order with different delays. I also added some parts to print out the received message and vector clock information at the time of the event. On my trials, every peer delivered the message to application in the correct order.

Code snippet 10: `postMessageToPeer` helper function

Figure 3: A message that is postponed since it does not satisfy the condition.

Figure 4: Sending of 9 consequent messages with random delays.

```

ec2-user@ip-172-31-95-231:~/Assignment2
1 Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg3
  TimeStamp: [0 0 0 0 3]
}
My current vector clock is [0 0 0 0 2]
Message is delivered
172.31.89.7:4444: msg3

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg8
  TimeStamp: [0 0 0 0 8]
}
My current vector clock is [0 0 0 0 3]
Message is postponed

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg5
  TimeStamp: [0 0 0 0 5]
}
My current vector clock is [0 0 0 0 3]
Message is postponed

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg4
  TimeStamp: [0 0 0 0 4]
}
My current vector clock is [0 0 0 0 3]
Message is delivered
172.31.89.7:4444: msg4
172.31.89.7:4444: msg5

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg7
  TimeStamp: [0 0 0 0 7]
}
My current vector clock is [0 0 0 0 5]
Message is postponed

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg5
  TimeStamp: [0 0 0 0 5]
}
My current vector clock is [0 0 0 0 5]
Message is delivered
172.31.89.7:4444: msg5

ec2-user@ip-172-31-95-231:~/Assignment2
2 Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg5
  TimeStamp: [0 0 0 0 5]
}
My current vector clock is [0 0 0 0 3]
Message is postponed

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg4
  TimeStamp: [0 0 0 0 4]
}
My current vector clock is [0 0 0 0 3]
Message is delivered
172.31.89.7:4444: msg4
172.31.89.7:4444: msg5

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg7
  TimeStamp: [0 0 0 0 7]
}
My current vector clock is [0 0 0 0 5]
Message is postponed

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg9
  TimeStamp: [0 0 0 0 9]
}
My current vector clock is [0 0 0 0 5]
Message is postponed

Received a message: {
  Sender ID: 172.31.89.7:4444
  Transcript: msg6
  TimeStamp: [0 0 0 0 6]
}
My current vector clock is [0 0 0 0 5]
Message is delivered
172.31.89.7:4444: msg6
172.31.89.7:4444: msg7
172.31.89.7:4444: msg8
172.31.89.7:4444: msg9

```

Figure 5: Peer 1 Delivering the messages in the correct order

References

- [1] Concurrent map and slice types in Go. Retrieved March 15, 2020, from <https://dnaeon.github.io/concurrent-maps-and-slices-in-go/>